

Know what to test before you ship.

We read your prompt changes and help you choose
which tests to run first.

Support-email example: V1 and the shorter V2 pass all five basic tests.

Production V1

100%

Candidate V2

100%

Looks safe to ship.

Qwen reads support emails. GPT-OSS chooses the tests.

Qwen3-30B · tested earlier

Reads each email using V1 and V2.

Example extraction

“Please refund order #1234.”

↓

Order: 1234 · Action: refund

Answers and scores saved before the demo.

GPT-OSS-120B · live

1. Read what changed in V1 → V2.
2. Choose an email test and explain why.
3. App loads that test's saved V1 and V2 scores.
4. See results → choose again or stop.

Can we break it?

TESTS NEEDED TO FIND THE PROBLEM

2, 2, 1 tests

Nebius planner fresh runs

3.5 tests expected

Random search · without replacement

Small synthetic demonstration; n=3 planner runs.

Hidden regression

Instructions hidden in the email

V1	40%
V2	0%

Δ -40 pp

Choose what to test before the next release

CUSTOMER	AI teams shipping prompt and model updates
PAIN	Choosing edge cases takes time. Running every test uses inference budget.
PRODUCT	Test selection inside the team's release workflow
BUSINESS	Team subscriptions + evaluation usage Private deployment for enterprise customers

Today: prompt regression testing.

Next: models, RAG systems and AI agents.