

Find what got worse before you ship what got better.

AI teams ship prompt and model updates constantly.

Standard evals can stay green while important user cases regress.

Production V1


100%

Candidate V2

100%

Looks safe to ship.

Break My Method actively decides what to test next.



The model chooses
where to look.
It never invents
the evidence.

Planning is live.
Evidence is executed, saved,
and scored against ground truth.

Can we break it?

AUDITED EVIDENCE

2, 2, 1 tests

Nebius planner fresh runs

3.5 tests expected

Random search · without replacement

Within 2 tests

Nebius: 3/3

Random: 33.3%

Small synthetic demonstration; n=3 planner runs.

Hidden regression

Instruction-like customer content

V1

40%

V2

0%

Δ -40 pp

From static eval suites to active AI testing

CUSTOMER	AI teams shipping prompt/model updates
PAIN	Manual regression-test selection consumes engineering time and inference budget
PRODUCT	Active testing layer on top of existing eval workflows
BUSINESS	Team subscription + evaluation usage + enterprise private deployment

Today: prompt regression testing.

Next: models, RAG systems and AI agents.