

Find what got worse before you ship what got better.

We shortened the instructions for an AI that reads support emails.

Both versions still passed all five basic tests.

Production V1


100%

Candidate V2

100%

Looks safe to ship.

The model picks a test. The app checks the result.



GPT-OSS-120B
chooses the test.
Qwen3-30B
supplied the answers.

Planning runs live.
Qwen outputs were saved earlier.

Can we break it?

TESTS NEEDED TO FIND THE PROBLEM

2, 2, 1 tests

Nebius planner fresh runs

3.5 tests expected

Random search · without replacement

Within 2 tests

Nebius: 3/3

Random: 33.3%

Small synthetic demonstration; n=3 planner runs.

Hidden regression

Instructions hidden in the email

V1	40%
V2	0%

Δ -40 pp

Choose what to test before the next release

CUSTOMER	AI teams shipping prompt and model updates
PAIN	Choosing edge cases takes time. Running every test uses inference budget.
PRODUCT	Test selection inside the team's release workflow
BUSINESS	Team subscriptions + evaluation usage Private deployment for enterprise customers

Today: prompt regression testing.

Next: models, RAG systems and AI agents.